

Basics of Data Analysis

Homework report

December 2023
Moscow, Russia

Author: Marwan Bouabidi
Supervisor: Prof. Mirkin Boris Grigorievich



NATIONAL RESEARCH
UNIVERSITY

Department of Computer Science

Contents

1	Introduction	3
2	Dataset	3
3	Analysis	4
3.1	Correlation coefficient	4
3.2	Principal Component Analysis (PCA)	6
3.3	K-means Clustering	9
3.3.1	K=4	11
3.3.2	K=7	12
3.3.3	Interpretation	13
3.4	Contingency tables	16
3.5	Bootstrap	18
4	Conclusion	20
5	Appendix	21
5.1	Python code	21

List of Figures

1	Scatter-plot of height and reach for different boxers with linear regression model	5
2	Errors of the Linear Regression model	6
3	SVD with z-scoring	7
4	SVD with range normalisation	7
5	PC visualisation with z-scoring	8
6	PC visualisation with range normalisation	8
7	Inertia values over 10 iterations for K=4	10
8	Inertia values over 10 iterations for K=7	10
9	KMeans clusters for K=4 visualised for features age and won with range normalisation.	11
10	KMeans clusters for K=4 visualised for features age and won with real data.	11
11	KMeans clusters for K=7 visualised for features age and won with range normalisation.	12
12	KMeans clusters for K=7 visualised for features age and won with real data.	12
13	Numbered Kmeans clusters for K=4 visualised for features age and won with real data.	15
14	Bins interval ranges.	16
15	Bootstrap results.	20

List of Tables

1	Contingency table	17
2	Frequency table	17
3	Conditional probability (or frequency) table	17
4	Quetelet index table	18

1 Introduction

In this small study, I will employ some powerful data analysis tools to reveal hidden patterns and relationships within the chosen dataset of boxing matches. The correlation coefficient will enable us to assess the connections between different variables. Principal Component Analysis (PCA) will assist in identifying the most significant features in the data. And K-means clustering will allow us to group matches based on shared characteristics, revealing hidden clusters of data. I will also use contingency tables to assess statistical relationship s between some features, and use the bootstrap method along with confidence intervals to compare means between 2 chosen clusters, and the grand mean with a cluster mean.

This analysis provides an interesting way of exploring the dynamics of the sport and seeing it from an unusual angle.

2 Dataset

This data set is a collection of boxing matches results, since I'm interested in the sport and thought it would be a good idea to apply some data analysis tools to this kind of information.

I removed rows with missing data (except for judges scores, as they are not always present) and was left with **7155** rows or objects in the resulting dataset. Initially, the features were:

- **age_A, age_B**: age of the boxer
- **height_A, height_B**: height (in cm)
- **reach_A, reach_B**: arms reach (in cm)
- **stance_A, stance_B**: type of stance, 'orthodox' or 'southpaw'
- **weight_A, weight_B**: weight (in lbs)
- **won_A, won_B**: number of wins
- **lost_A, lost_B**: number of losses
- **drawn_A, drawn_B**: number of draws
- **kos_A, kos_B**: number of KOs
- **result**: the match result. 'win_A', 'win_B', or 'draw'
- **decision**: match conclusion. 'UD' (unanimous decision), 'SD' (split draw), 'MD' (majority decision), 'PTS' (points), 'KO', 'TKO' (technical KO), 'TD', 'RTD' (retired), 'DQ' (disqualification)

Please see this link for abbreviations:

<https://www.boxingbase.com/boxing-results-explained/>

- `judge1_A`, `judge1_B`: scores of judge 1
- `judge2_A`, `judge2_B`: scores of judge 2
- `judge3_A`, `judge3_B`: scores of judge 3

However, I was advised to consider the matches from the point of view of Boxer A only. So the features became something like: `age_A`, `age_diff`, `height_A`, `height_diff`, `reach_A`, `reach_diff`... Where:

$$\text{feature_diff} = \text{feature_A} - \text{feature_B}$$

And so, the new list of features is:

```
['age', 'age_diff', 'height', 'height_diff',
 'reach', 'reach_diff', 'stance', 'stance_diff',
 'weight', 'weight_diff', 'won', 'won_diff',
 'lost', 'lost_diff', 'drawn', 'drawn_diff', 'kos',
 'kos_diff', 'result', 'decision', 'judge1',
 'judge1_diff', 'judge2', 'judge2_diff', 'judge3',
 'judge3_diff']
```

So, if the difference is **positive**, the feature for the boxer we're considering is higher than that of his opponent, and vice-versa for the negative differences.

Source: <https://www.kaggle.com/datasets/mexwell/boxing-matches>

3 Analysis

3.1 Correlation coefficient

We need firstly to find 2 features with a linear-like relationship, and that is achieved by examining scatter-plots of feature pairs. I have chosen to examine the classic features of **height** and **reach** of the boxer, as their scatter-plot shows a linear pattern, with only a couple of points falling far from the rest, also known as **outliers**, which could greatly affect the linear regression model, in particular by inflating the correlation coefficient ρ . We can, of course, rule some of these points out as errors in measurements, because, for example, a reach of more than 4 meters doesn't seem humanly possible. Also, it is not irrational to assume some form of linear relationship between these 2 features, because based on everyday observations it seems they might be proportional (although, to my knowledge, there is no clear scientific evidence for this).

The scatter-plot below shows a linear-like pattern of points suitable enough to build a **Linear Regression** model, so let us examine this further.

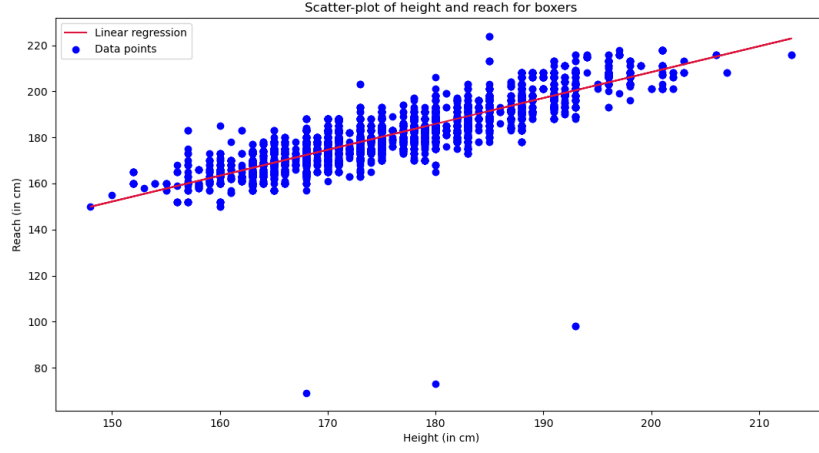


Figure 1: Scatter-plot of height and reach for different boxers with linear regression model

For 2 features \mathbf{X} and \mathbf{Y} , correlation coefficients are given by the following formulas:

$$\rho = \frac{\sum_{i=1}^N (x_i - \bar{X})(y_i - \bar{Y})/N}{\sigma(x)\sigma(y)}$$

$$a = \rho \frac{\sigma(y)}{\sigma(x)}$$

$$b = \bar{y} - a * \bar{x}$$

$$\text{Such that: } Y = a * X + b$$

Where σ is the standard deviation, ρ the **correlation coefficient**, Y representing the reach, and X the height.

In python, I calculated 2 versions of ρ , one with the **scipy.stats** library, and another using the formula. Rounded to two decimals, they're equal.

For these chosen features, I obtained:

$$\rho = 0.88$$

$$a = 1.12$$

$$b = -16.45$$

$$\text{Therefore: } Y = 1.12 * X - 16.45$$

$$\text{The determinancy coefficient: } \rho^2 = 0.77 = 77\%$$

The correlation coefficient measures the extent of linearity between X and Y , while the determinancy coefficient represents the percentage of variance of Y

that is explained by linear regression (of Y over X).
We compute two errors to verify the regression model:

- **Data Analysis error:** $\frac{|Y_{pred}-Y_{real}|}{|Y_{real}|} * 100$
- **Machine Learning error:** $\frac{|Y_{pred}-Y_{real}|}{|Y_{pred}|} * 100$

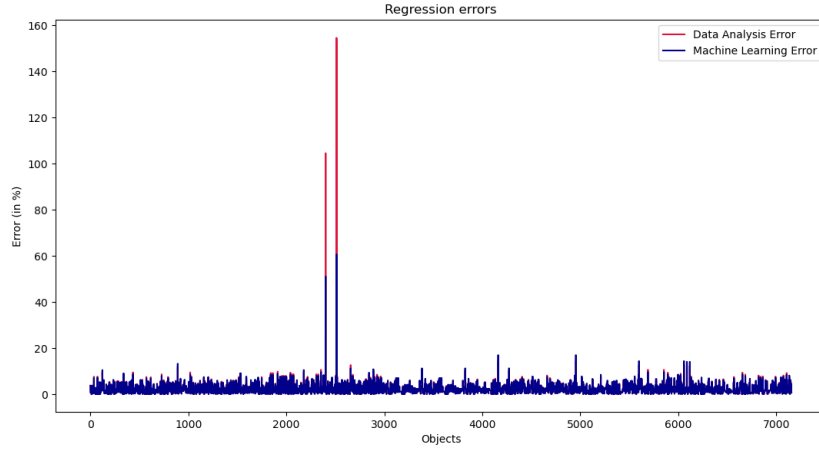


Figure 2: Errors of the Linear Regression model

Both errors are almost always below the 20% line, except for a few spikes, so we can conclude that the linear regression is a decent model to represent the relationship between these 2 features. The determinancy coefficient is 77%, leaving around 23% error margin.

3.2 Principal Component Analysis (PCA)

We assume our data are arranged in a matrix \mathbf{X} , where the lines represent the objects, and the columns the features.

The features below were chosen because I believe they influence the outcome of the match the most.

`['age', 'age_diff', 'won', 'won_diff', 'lost', 'lost_diff']`

We start by standardizing the data (for every column), and for that we'll be using 2 methods:

- **Z-scoring:** $Y = \frac{X - \bar{X}}{\sigma_X}$
- **Range normalisation:** $Y = \frac{X - \bar{X}}{max - min}$

After, we apply SVD on the data matrix X for both cases (we'll note them X_z and X_r):

$$Z, \mu, C = SVD(X)$$

Where: $X = Z * \mu * C^T$

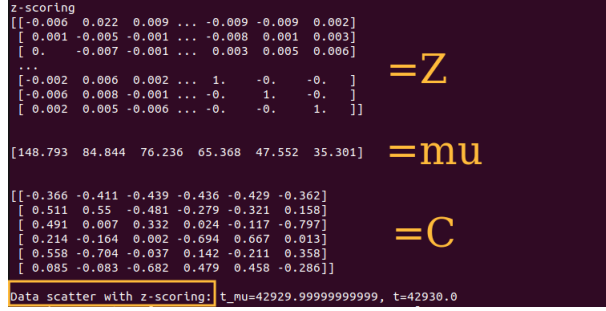


Figure 3: SVD with z-scoring

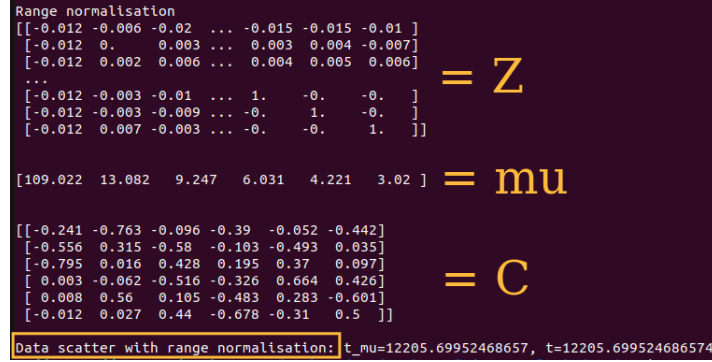


Figure 4: SVD with range normalisation

We notice from the SVD results and data scatter (which was calculated twice, with Y and μ) that:

Putting the square of the singular values in μ in percentage, we can see the contribution of each of the principal components to the data scatter:

$$\begin{aligned} \mu_{2_z} &= [51.57, 16.77, 13.54, 9.95, 5.27, 2.9] \\ \mu_{2_r} &= [54.59, 20.59, 11.23, 7.45, 4.18, 1.95] \end{aligned}$$

Visualisation:

PCA helps us project our data objects on a 2D plane. For both cases, we visualise the first 2 components with the most contribution, which are the first 2 columns of Z : Z_1 and Z_2 . And, depending on the values for C 's columns, we

can consider C_i (and Z_i) or $-C_i$ (and $-Z_i$) because they are all eigenvectors of $X^T X$, since we have: $XC = \mu Z$ and $X^T Z = \mu C$ for a singular value μ . So, for visualization:

- For z-scoring: Z_1 and $-Z_2$
- For range normalisation: $-Z_1$ and $-Z_2$

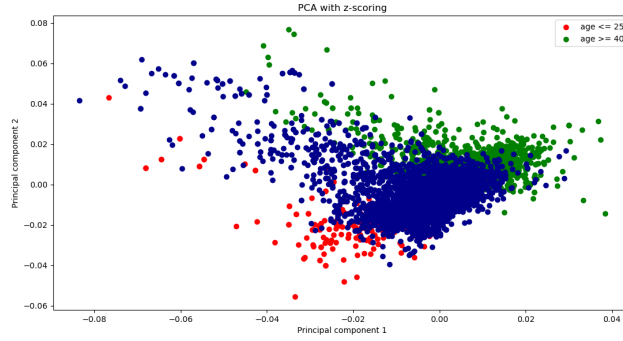


Figure 5: PC visualisation with z-scoring

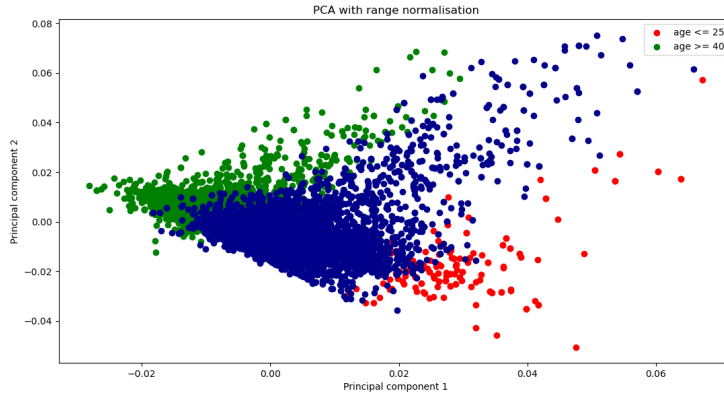


Figure 6: PC visualisation with range normalisation

By examining the corresponding scatters, we can see some clusters forming after age filters have been applied. The results of both standardisation methods look similar, since their respective principal components have similar contributions to the data scatter. But visually, it seems the data is more separated and clearer for interpretation with range normalisation.

Hidden factor:

To find a hidden factor, we start by normalising the data to get values between 0 and 100 using ranking normalisation:

$$Y = \frac{X - \min}{\max - \min} * 100$$

Then, similarly, we apply SVD to get the principal components, we choose the most dominant, we get rid of negative values if they dominate, and finally the hidden factor is computed as:

$$\alpha = \frac{1}{\text{sum}(C)}$$

In our case, the value rounded to 2 decimals is $\alpha = 0.8 = 80\%$.

3.3 K-means Clustering

K-Means is an unsupervised Machine Learning algorithm used to partition or divide data objects into K groups or clusters, based on euclidean distances of these points to the K cluster centers.

For our dataset, we will apply this algorithm twice: one for **K=4**, and another for **K=7**. For each case we will run the algorithm 10 times, with random initialisations, and choose the best based on **inertia** value:

$$D(S, c) = \sum_{p=1}^K \sum_{i \in S_p} d(i, c_p) = \sum_{p=1}^K \sum_{i \in S_p} \sum_{v=1}^V (y_{iv} - c_{kv})^2$$

I chose the below features:

```
features = ['age', 'age_diff', 'reach', 'reach_diff', 'won',
            'won_diff', 'kos', 'kos_diff']
```

These features, in my personal opinion, affect the most the outcome of the match. After standardizing with range normalisation (in%) the data and running the algorithm multiple times, we obtain a list of inertia values. We choose the iteration with minimum inertia value.

Something to note here: in python, the **KMeans** library specifically, there is an option `n_init` that controls how many times the algorithm will be executing, and then Python will return automatically the best iteration based on inertia values. It was set to 10 in my code, and I also executed KMeans 10 times manually, for the sake of showing results, so KMeans was actually executed a **100** times.

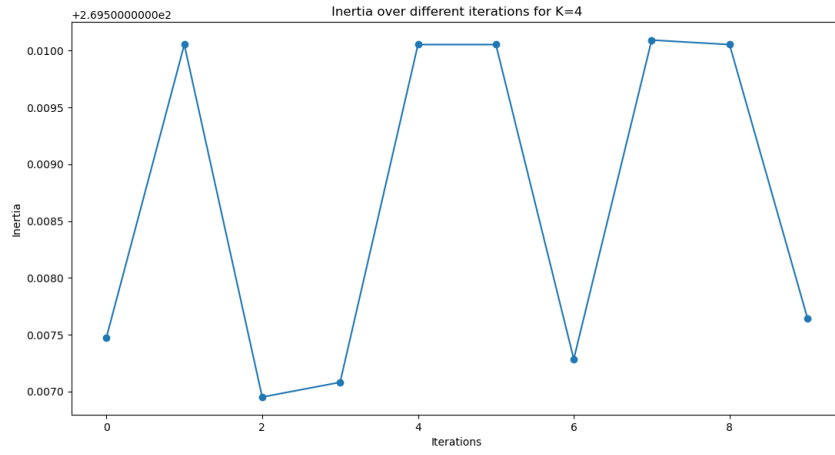


Figure 7: Inertia values over 10 iterations for K=4

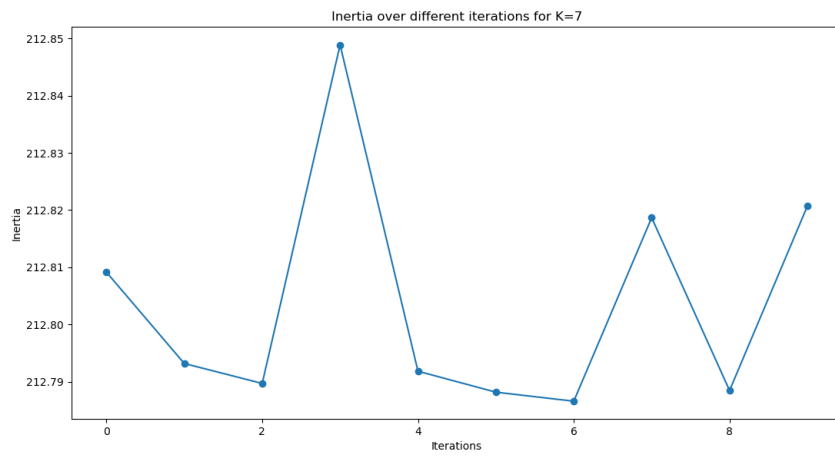


Figure 8: Inertia values over 10 iterations for K=7

3.3.1 K=4

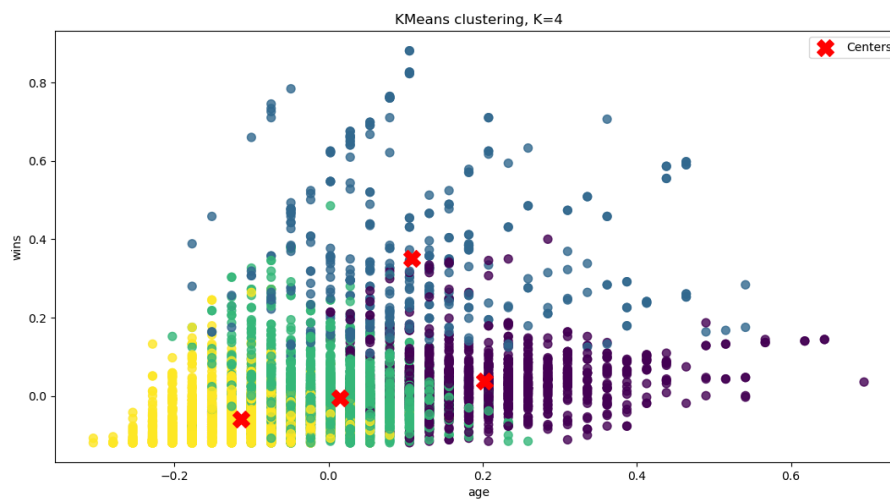


Figure 9: KMeans clusters for K=4 visualised for features **age** and **won** with range normalisation.

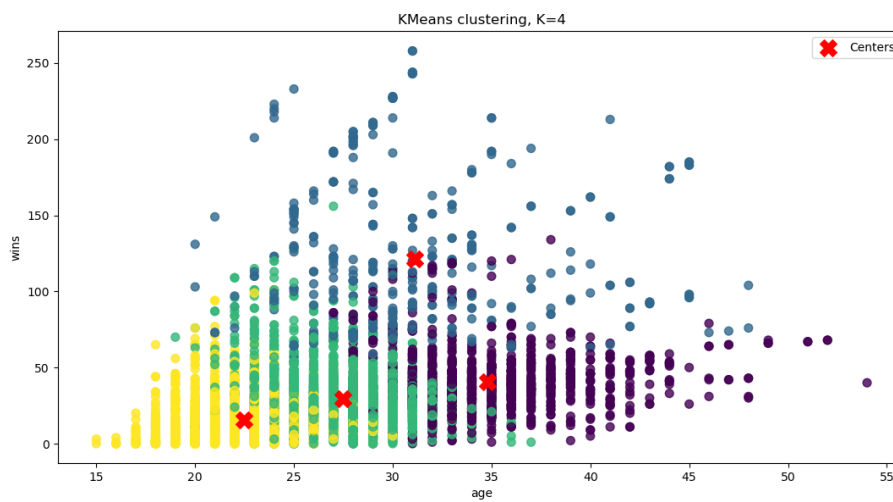


Figure 10: KMeans clusters for K=4 visualised for features **age** and **won** with real data.

3.3.2 K=7

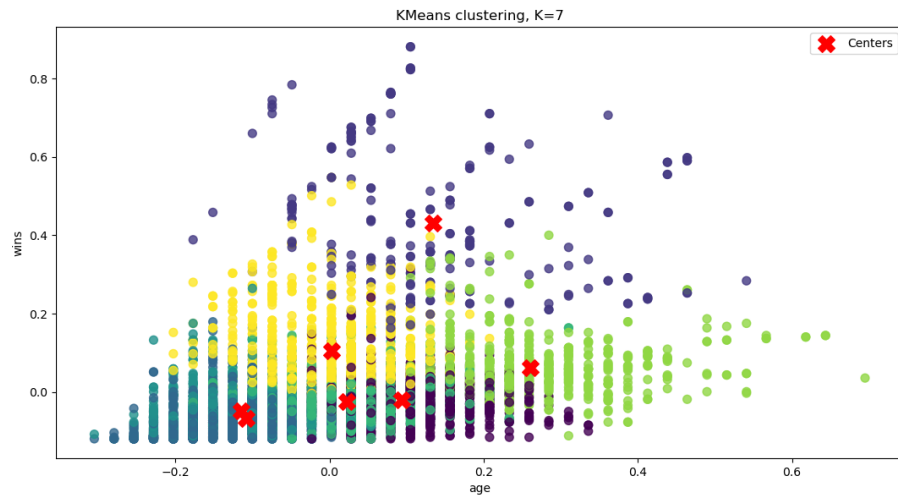


Figure 11: KMeans clusters for K=7 visualised for features **age** and **wins** with range normalisation.

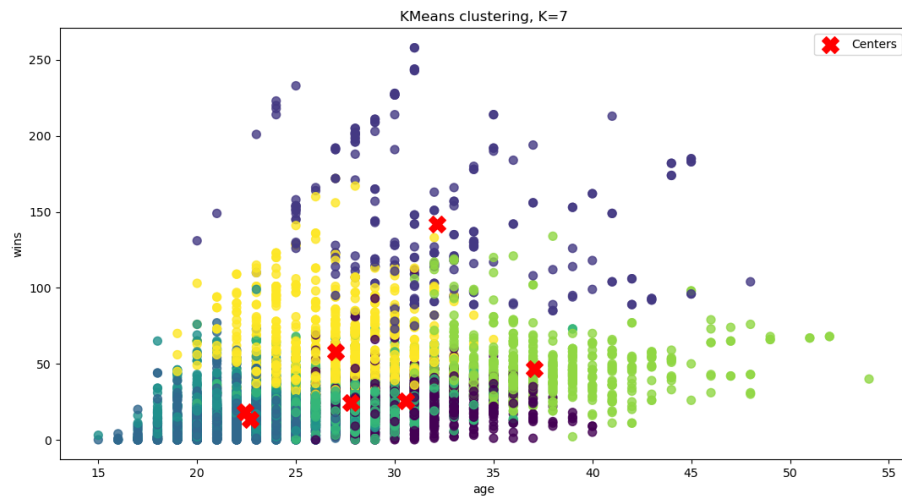


Figure 12: KMeans clusters for K=7 visualised for features **age** and **wins** with real data.

3.3.3 Interpretation

By comparing the local means of clusters to the global mean for every chosen feature, I obtained the following results for K=4:

```
#Global means:
{'age': 26.91, 'age_diff': -1.41, 'reach': 181.46, 'reach_diff': 0.88,
'won': 30.75, 'won_diff': 3.26, 'kos': 18.23, 'kos_diff': 2.35}

#Means for cluster number 1 (size 1134) for K=4:
{'age': 34.8, 'age_diff': 6.34, 'reach': 185.4, 'reach_diff': 0.61,
'won': 40.48, 'won_diff': 14.46, 'kos': 25.38, 'kos_diff': 8.67}
#Relative differences to global means (%):
[22.67, 122.24, 2.13, 44.26, 24.04, 77.46, 28.17, 72.9]
Average (%): 49.23

#Means for cluster number 2 (size 376) for K=4:
{'age': 31.11, 'age_diff': 4.58, 'reach': 180.57, 'reach_diff': -2.39,
'won': 121.15, 'won_diff': 72.61, 'kos': 59.76, 'kos_diff': 36.6}
#Relative differences to global means (%):
[13.5, 130.79, 0.49, 136.82, 74.62, 95.51, 69.49, 93.58]
Average (%): 76.85

#Means for cluster number 3 (size 2871) for K=4:
{'age': 27.5, 'age_diff': -0.62, 'reach': 182.12, 'reach_diff': 0.82,
'won': 29.47, 'won_diff': 1.96, 'kos': 18.18, 'kos_diff': 1.96}
#Relative differences to global means (%):
[2.15, 127.42, 0.36, 7.32, 4.34, 66.33, 0.28, 19.9]
Average (%): 28.51

#Means for cluster number 4 (size 2774) for K=4:
{'age': 22.5, 'age_diff': -6.21, 'reach': 179.29, 'reach_diff': 1.49,
'won': 15.84, 'won_diff': -9.37, 'kos': 9.72, 'kos_diff': -4.46}
#Relative differences to global means (%):
[19.6, 77.29, 1.21, 40.94, 94.13, 134.79, 87.55, 152.69]
Average (%): 76.03
```

And for K=7, I obtained the following results:

```
#Global means:
{'age': 26.91, 'age_diff': -1.41, 'reach': 181.46, 'reach_diff': 0.88,
'won': 30.75, 'won_diff': 3.26, 'kos': 18.23, 'kos_diff': 2.35}

#Means for cluster number 1 (size 1149) for K=7:
{'age': 30.56, 'age_diff': 4.3, 'reach': 177.48, 'reach_diff': -1.58,
'won': 25.9, 'won_diff': 3.86, 'kos': 15.28, 'kos_diff': 1.13}
```

```

#Relative differences to global means (%):
[11.94, 132.79, 2.24, 155.7, 18.73, 15.54, 19.31, 107.96]
Average (%): 58.03

#Means for cluster number 2 (size 232) for K=7:
{'age': 32.16, 'age_diff': 5.66, 'reach': 180.43, 'reach_diff': -3.19,
'won': 141.76, 'won_diff': 95.12, 'kos': 66.2, 'kos_diff': 45.19}
#Relative differences to global means (%):
[16.32, 124.91, 0.57, 127.59, 78.31, 96.57, 72.46, 94.8]
Average (%): 76.44

#Means for cluster number 3 (size 1623) for K=7:
{'age': 22.7, 'age_diff': -1.34, 'reach': 175.15, 'reach_diff': 0.42,
'won': 13.47, 'won_diff': -0.95, 'kos': 7.89, 'kos_diff': -0.03}
#Relative differences to global means (%):
[18.55, 5.22, 3.6, 109.52, 128.29, 443.16, 131.05, 7933.33]
Average (%): 1096.59

#Means for cluster number 4 (size 1312) for K=7:
{'age': 22.42, 'age_diff': -9.7, 'reach': 180.31, 'reach_diff': 1.67,
'won': 18.59, 'won_diff': -15.98, 'kos': 11.6, 'kos_diff': -7.78}
#Relative differences to global means (%):
[20.03, 85.46, 0.64, 47.31, 65.41, 120.4, 57.16, 130.21]
Average (%): 65.83

#Means for cluster number 5 (size 1403) for K=7:
{'age': 27.76, 'age_diff': -3.78, 'reach': 191.0, 'reach_diff': 4.19,
'won': 24.58, 'won_diff': -5.84, 'kos': 16.51, 'kos_diff': -2.92}
#Relative differences to global means (%):
[3.06, 62.7, 4.99, 79.0, 25.1, 155.82, 10.42, 180.48]
Average (%): 65.2

#Means for cluster number 6 (size 593) for K=7:
{'age': 37.07, 'age_diff': 7.63, 'reach': 190.29, 'reach_diff': 2.19,
'won': 46.79, 'won_diff': 19.86, 'kos': 30.75, 'kos_diff': 13.2}
#Relative differences to global means (%):
[27.41, 118.48, 4.64, 59.82, 34.28, 83.59, 40.72, 82.2]
Average (%): 56.39

#Means for cluster number 7 (size 843) for K=7:
{'age': 27.0, 'age_diff': -0.82, 'reach': 179.02, 'reach_diff': -1.42,
'won': 57.96, 'won_diff': 18.68, 'kos': 33.32, 'kos_diff': 13.74}
#Relative differences to global means (%):
[0.33, 71.95, 1.36, 161.97, 46.95, 82.55, 45.29, 82.9]
Average (%): 61.66

```

The features **age** and **reach** seem to present the lowest relative errors in both partitions, while **age_diff**, **won_diff**, **reach_diff** (occasionally) and **kos_diff** seem to have the highest values. Partitions for **K=4** present lower relative errors to the global mean, and also they look more separate graphically, therefore my choice is for **K=4** as the most suited for interpretation of the 2.

Looking at the results for the **K=4** partition, we can see 4 categories of boxers:

1. First, those with the highest wins (average of 121.15) spanning all over the age range, represented by **cluster 2 with 376 objects**, but mainly scattered around their mean age of **31.11** (31 often considered the prime age of boxers).
2. Those with the lower wins (looking at some of them, they have around 50 wins, which is a great number in boxing, but here we don't take into account some other factors such as the quality of those wins, and so the algorithm is considering only this number, along with the other features) divided into 3 age groups themselves: young (mean **22.5**) represented by **cluster 4 with 2774 objects**, prime age (mean **27.5**) represented by **cluster 3 with 2871 objects**, and old (mean **34.8**) represented by **cluster 1 with 1134 objects**.

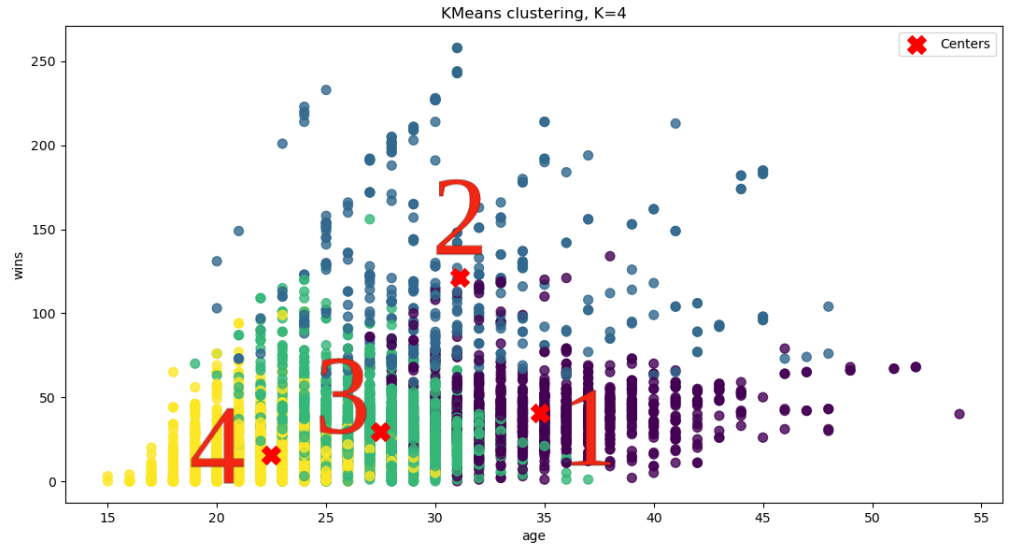


Figure 13: Numbered Kmeans clusters for **K=4** visualised for features **age** and **won** with real data.

It seems, those with the highest number of wins are mainly those old and active enough to accumulate these wins with time, but nevertheless we can see some talents with high wins at relatively young age.

Now, looking at the computed means we can see some interesting observations:

- Boxers with high wins (cluster 2) tend to choose younger opponents (with an average age of 31 in this cluster, one could say that boxers 4.58 years younger are considered inexperienced), with much less wins and KOs (average wins difference is 72.6, and for KOs it's 36.6). So it seems the opponents are carefully chosen.
- Boxers with lower wins and old age (cluster 1): they choose younger opponents in their prime years, with less wins in general (14.46 average wins difference and 8.6 KOs difference with 40.48 average win record for this cluster, which seems acceptable)
- Boxers with lower wins and prime age (cluster 3): choosing tough opponents in their own age and win record level to keep their ranks or rise even more.
- Boxers with lower wins and young age (cluster 4): aspiring young boxers choosing older more experienced opponents (6.2 years older on average), with more wins (-9.37 win difference, -4.46 KOs difference) to rise in the ranks.

3.4 Contingency tables

In order to display the multivariate frequency distribution of the variables, we form a contingency table with 4 bins for the two following chosen nominal features: **won** and **won_diff** to analyse the choice of opponents based on win records. The data has been discretised into **4 bins** with the **quantile** strategy in Python, which means all bins in each feature have a similar number of points. The results are as follows:

```
For feature won, we have:
Bin number 1: [0, 13]
Bin number 2: [14, 23]
Bin number 3: [24, 36]
Bin number 4: [37, 258]
For feature won_diff, we have:
Bin number 1: [-205, -8]
Bin number 2: [-7, 0]
Bin number 3: [1, 9]
Bin number 4: [10, 211]
```

Figure 14: Bins interval ranges.

won_diff						
		G1	G2	G3	G4	Total
won	H1	379	694	596	13	1682
	H2	616	555	535	106	1812
	H3	410	343	519	543	1815
	H4	260	108	226	1252	1846
	Total	1665	1700	1876	1914	7155

Table 1: Contingency table

We deduce from the previous table, the table of relative frequencies by dividing every element by the total number of objects:

won_diff						
		G1	G2	G3	G4	Total
won	H1	0.05	0.1	0.08	0	0.24
	H2	0.09	0.08	0.07	0.01	0.25
	H3	0.06	0.05	0.07	0.08	0.25
	H4	0.04	0.02	0.03	0.17	0.26
	Total	0.23	0.24	0.26	0.27	1

Table 2: Frequency table

The conditional probability table is given by dividing every cell in the contingency table by the sum of elements of its corresponding row:

won_diff						
		G1	G2	G3	G4	Total
won	H1	0.23	0.41	0.35	0.01	1
	H2	0.34	0.31	0.3	0.06	1
	H3	0.23	0.19	0.29	0.3	1
	H4	0.14	0.06	0.12	0.68	1

Table 3: Conditional probability (or frequency) table

We can now compute the **Quetelet index table** defined as follows:

$$q_{ij} = q(j/i) = \frac{p(j/i) - p_j}{p_j} = \frac{NN_{ij}}{N_i N_j} - 1$$

Where $p(j/i) = p_{ij}/p_i$ is the (i, j) element from the conditional frequency table, and $p_j = N_j/N$ the probability for a column j .

		won_diff			
		G1	G2	G3	G4
won	H1	-0.03	0.74	0.35	-0.97
	H2	0.46	0.29	0.13	-0.78
	H3	-0.03	-0.2	0.09	0.12
	H4	-0.39	-0.75	-0.53	1.54

Table 4: Quetelet index table

From here we can calculate:

$$\text{Average Quetelet Index: } q_{avg} = \sum_{i,j} p_{ij} * q_{ij} = 0.39$$

$$\text{Chi-squared: } \Phi^2 = \sum_{i,j} \frac{(p_{ij} - p_i p_j)^2}{p_i p_j} = 0.39$$

$$\text{Number of degrees of freedom: } (N_{rows} - 1) * (N_{cols} - 1) = 9$$

$$\text{Pearson's Chi-squared: } N * \Phi^2 = 2771.9; \text{ N the number of objects}$$

Chi-squared or Φ^2 measures deviation from statistical independence. The average Quetelet index is interpreted as: the knowledge of a category of one feature (i) increases knowledge of (j) category by $100 * q_{avg}\% = 39\%$. So, knowing a boxer's win records gives a decent chance of knowing his opponent's win record, which makes sense.

According to chi-squared tables for 9 DF:

- For 95% confidence level we need $19.92/0.39 = 43.38$ observations, so **44** minimum.
- For 99% confidence level we need $19.02/0.39 = 48.77$ observations, so **49** minimum.

3.5 Bootstrap

Bootstrap is a statistical resampling technique used to estimate the sampling distribution of a statistic. Its main idea is to simulate multiple datasets that resemble the original data.

From the clustering section, we choose 2 clusters with the lowest average relative errors from the **K=4** partition. We perform bootstrap on these clusters with the nominal feature **won** (representing the number of wins).

My dataset has **N=7155 objects**. I generated a matrix of indices of size **5000*N**, where every line has N random indices between **0 and N-1**, then I calculated the averages for every line (for the chosen feature 'won') and stored it in a vector **mx**. To find a **95% confidence interval** for the **grand mean=30.75**, I used two method:

- **Pivotal method:**

$$Ic = [\text{mean}(mx) - 1.96 * \sigma(mx), \text{mean}(mx) + 1.96 * \sigma(mx)] = [30.07, 31.42]$$

- **Non-pivotal method:** we sort the values of mx into mxs , then we remove 2.5% from each side of mxs , so we get:

$$Ic = [mxs[125], mxs[4874]] = [30.08, 31.42]$$

To compare the means between clusters, or the mean of a cluster with the grand mean, we use a similar methodology. First, we construct 2 vectors of 5000 averages for each of the clusters. This is done as follows: for a chosen cluster CL_i , we through all the 5000 tries generated earlier, and for each try, containing N random indices, we select the indices that are in the cluster CL_i . Then we average these values, and we obtain one average for 1 try. We repeat this 5000 times to get the wanted vector. Let us call our chosen cluster, CL_1 and CL_2 , and their respective vectors of averages $mx1$ and $mx2$. We can compare the means of the clusters CL_1 and CL_2 by considering the vector $m12 = mx1 - mx2$, and then finding, as we previously did, a 95% confidence interval for the mean of this difference. If 0 is in that interval, we can accept the hypothesis that $\text{mean}(CL_1) = \text{mean}(CL_2)$. Similarly, we compare the mean of one of the clusters to the grand mean (the latter represented by the vector mx).

After applying this, I got the following results:

- **Pivotal method:**

$$Ic_{12} = [-12.03, -9.62]$$

$$Ic_{1g} = [0.33, 1.81]$$

$$Ic_{2g} = [-10.84, -8.68]$$

- **Non-pivotal method:**

$$Ic_{12} = [-12.02, -9.6]$$

$$Ic_{1g} = [0.34, 1.8]$$

$$Ic_{2g} = [-10.8, -8.65]$$

So we can conclude that the means of the clusters and the grand mean are all different. None of the equality hypotheses were accepted since 0 is not in any of the confidence intervals found with our samples. The closest hypothesis to being true was: $\text{mean}(CL_1) = \text{grand mean}$.

```
Bootstrap method:
Confidence interval for the grand mean:
Pivotal Ic=[30.07344669472215, 31.420212620442072]
Non-pivotal Ic=[30.07756813417191, 31.41928721174004]
Grand mean: 30.749685534591194
Comparing cl1 and cl2 means:
Pivotal Ic=[-12.025649750535875, -9.62108178945205]
Non-pivotal Ic=[-12.025038132843271, -9.608605371882437]
Comparing cl1 and grand mean:
Pivotal Ic=[0.3269598652708896, 1.8094592653227726]
Non-pivotal Ic=[0.3431111051148683, 1.801782950292914]
Comparing cl2 and grand mean:
Pivotal Ic=[-10.83507999736562, -8.675232412028643]
Non-pivotal Ic=[-10.804203526650433, -8.645049329478887]
```

Figure 15: Bootstrap results.

4 Conclusion

By analysing this boxing dataset, I was able to visualise and analyse some of the effects of important factors in the sport such as win record, losses, number of KOs, age, height, reach, etc... relying on different data analysis tools.

I was also able to notice some interesting but known trends in the data when it comes to the choice of opponents for different groups of boxers.

5 Appendix

5.1 Python code

```
1  #Data preprocessing: Line 20
2  #Correlation: Line 48
3  #PCA: Line 92
4  #KMeans Clustering: Line 188
5  #Contingency Table: Line 310
6  #Bootstrap: Line 368
7  #Main program: Line 462
8
9  import pandas as pd
10 import matplotlib.pyplot as plt
11 import numpy as np
12 from scipy.stats import pearsonr
13 from sklearn.cluster import KMeans
14 from sklearn.preprocessing import KBinsDiscretizer
15
16 #Cluster indices for bootstrap
17 cl1_indices = []
18 cl2_indices = []
19
20 def preprocessing():
21     try:
22         path = "Datasets/Boxing/"
23         file_name = 'boxing_matches.csv'
24         df = pd.read_csv(path+file_name)
25         ignore_columns = ['judge1_A', 'judge1_B', 'judge2_A', 'judge2_B', 'judge3_A',
26                             'judge3_B']
27         df.dropna(subset=df.columns.difference(ignore_columns), inplace=True)
28         df.reset_index(drop=True, inplace=True)
29         df['result'] = df['result'].replace('win_A', 'win').replace('win_B', 'loss')
30         #Categories
31         #print(df['result'].value_counts())
32         cols = df.columns.tolist()
33         for i,col in enumerate(cols):
34             if ('_B' in col) and (df[col].dtype.kind in 'iufc'):
35                 df[col] = df[cols[i-1]] - df[col]
36         new_cols = {col:col.replace('_A','').replace('_B','_diff') for col in cols}
37         df.rename(columns=new_cols, inplace=True)
38         #Reach of more than 4m isnt acceptable, drop rows
39         df.drop(labels=[2805,2963,2893], axis=0, inplace=True)
40         df.reset_index(drop=True, inplace=True)
41         df.to_csv(path+file_name.replace('.csv', '_updated.csv'), index=False)
42         return df
```

```

43     except Exception as e:
44         print(r'Unable to read data: {}'.format(str(e)))
45         return pd.DataFrame()
46
47
48 def correlation(df):
49     try:
50         print('Correlation:')
51         #plt.matshow(df.corr(numeric_only=True))
52         #plt.show()
53         #pd.plotting.scatter_matrix(df, alpha=0.5, figsize=(6, 6), diagonal='hist')
54         X = df['height'].to_numpy()
55         Y = df['reach'].to_numpy()
56         #Reach of more than 4m is irrational, drop those rows
57         #Rho from library
58         rho = pearsonr(X,Y).statistic
59         #Rho built with formula
60         rho_ = sum([(X[i]-np.mean(X))*(Y[i]-np.mean(Y)) for i in range(len(X))]) /
61             (len(X)*np.std(X)*np.std(Y))
62
63         print('Library rho = {}, calculated rho={}'.format(rho,rho_))
64         a = rho*np.std(Y)/np.std(X)
65         b = np.mean(Y) - a*np.mean(X)
66         print('a={}, b={}'.format(a,b))
67         print('Correlation coeff: {}, determinancy coeff: {}'.format(rho,rho**2))
68         #Scatter-plot
69         plt.plot(X, [a*x+b for x in X], color='crimson', label='Linear regression')
70         plt.scatter(X, Y, color='blue', label='Data points')
71         plt.xlabel('Height (in cm)')
72         plt.ylabel('Reach (in cm)')
73         plt.legend(loc='upper left')
74         plt.title('Scatter-plot of height and reach for boxers')
75         plt.show()
76         #Compute errors
77         labels = [i for i in range(len(X))]
78         error_DS = [100*abs(a*x+b-Y[i])/abs(Y[i]) for i, x in enumerate(X)]
79         error_ML = [100*abs(a*x+b-Y[i])/abs(a*x+b) for i, x in enumerate(X)]
80         plt.plot(labels, error_DS, color='crimson', label='Data Analysis Error')
81         plt.plot(labels, error_ML, color='darkblue', label='Machine Learning Error')
82         plt.xlabel('Objects')
83         plt.ylabel('Error (in %)')
84         plt.legend(loc='upper right')
85         plt.title('Regression errors')
86         plt.show()
87     except Exception as e:
88         print(r'Unable to perform linear regression: {}'.format(str(e)))

```

```

89
90
91
92 def PCA(df):
93     try:
94         print('\n\nPrincipal component analysis:')
95         features = ['age', 'age_diff', 'won', 'won_diff', 'lost', 'lost_diff']
96         PCA_df_z = pd.DataFrame()
97         PCA_df_r = pd.DataFrame()
98         HiddenFactor_df = pd.DataFrame()
99
100        for feature in features:
101            #Standardization: z-scoring
102            sigma = np.std(df[feature])
103            mean = np.mean(df[feature])
104            max = df[feature].max()
105            min = df[feature].min()
106            PCA_df_z[feature] = df[feature].apply(lambda x: (x-mean)/sigma)
107            PCA_df_r[feature] = df[feature].apply(lambda x: (x-mean)/(max-min))
108            HiddenFactor_df[feature] = df[feature].apply(lambda x: 100*(x-min)/(max-min))
109
110
111        #Visualisation indices
112        df_inf = df.loc[df['age']>=40]
113        df_sup = df.loc[df['age']<=25]
114        #print(df_inf.index.values.tolist())
115        inf_indices = df_inf.index.values.tolist()
116        sup_indices = df_sup.index.values.tolist()
117        med_indices = [i for i in range(len(df)) if i not in inf_indices + sup_indices]
118
119        #####
120        ##### PCA z-scoring #####
121        #####
122
123        print('z-scoring')
124        X_z = PCA_df_z.to_numpy()
125        Z_z, mu_z, C_z = np.linalg.svd(X_z, full_matrices=True)
126        print(np.around(Z_z, decimals=3))
127        print('\n')
128        print(np.around(mu_z, decimals=3))
129        print('\n')
130        print(np.around(C_z, decimals=3))
131        #Data scatter
132        tm_z = sum([mu_i**2 for mu_i in mu_z])
133        t_z = np.sum(X_z**2)
134        print('\nData scatter with z-scoring: t_mu={}, t_z={}'.format(tm_z, t_z))

```



```

135     print('mu_z in percentage', [round(100*mu_i**2/t_z, 2) for mu_i in mu_z])
136     Z_1 = Z_z[:,0]
137     Z_2 = -Z_z[:,1]
138     plt.scatter(Z_1[inf_indices], Z_2[inf_indices], color='red', label="age <= 25")
139     plt.scatter(Z_1[sup_indices], Z_2[sup_indices], color='green', label="age >= 40")
140     plt.scatter(Z_1[med_indices], Z_2[med_indices], color='darkblue')
141     plt.xlabel('Principal component 1')
142     plt.ylabel('Principal component 2')
143     plt.legend(loc='upper right')
144     plt.title('PCA with z-scoring')
145     plt.show()
146
147     #####
148     ##### PCA Range Normalisation #####
149     #####
150
151     print('\n\nRange normalisation')
152     X_r = PCA_df_r.to_numpy()
153     Z_r, mu_r, C_r = np.linalg.svd(X_r, full_matrices=True)
154     print(np.around(Z_r, decimals=3))
155     print('\n')
156     print(np.around(mu_r, decimals=3))
157     print('\n')
158     print(np.around(C_r, decimals=3))
159     #Data scatter
160     tm_r = sum([mu_i**2 for mu_i in mu_r])
161     t_r = np.sum(X_r**2)
162     print('\nData scatter with range normalisation: t_mu={}, t={} '.format(tm_r, t_r))
163     print('mu_r in percentage', [round(100*mu_i**2/t_r, 2) for mu_i in mu_r])
164     Z_1 = -Z_r[:,0]
165     Z_2 = -Z_r[:,1]
166     plt.scatter(Z_1[inf_indices], Z_2[inf_indices], color='red', label="age <= 25")
167     plt.scatter(Z_1[sup_indices], Z_2[sup_indices], color='green', label="age >= 40")
168     plt.scatter(Z_1[med_indices], Z_2[med_indices], color='darkblue')
169     plt.xlabel('Principal component 1')
170     plt.ylabel('Principal component 2')
171     plt.legend(loc='upper right')
172     plt.title('PCA with range normalisation')
173     plt.show()
174
175     #####
176     ##### Hidden Factor #####
177     #####
178     X_h = HiddenFactor_df.to_numpy()
179     Z_h, mu_h, C_h = np.linalg.svd(X_h, full_matrices=True)
180     print(np.around(C_h, decimals=3))

```

```

181         C_1 = -C_h[:,0]
182         alpha = 1/sum(C_1)
183         print(alpha)
184     except Exception as e:
185         print(r'Unable to apply PCA: {}'.format(str(e)))
186
187
188 def KMeansClustering(df):
189     try:
190         print('\n\nKMeans Clustering:')
191         features = ['age', 'age_diff', 'reach', 'reach_diff', 'won', 'won_diff', 'kos',
192                     'kos_diff']
193         KMeans_df = pd.DataFrame()
194         global_means = {}
195         #Standardisation of the data
196         for feature in features:
197             #Standardization: z-scoring
198             sigma = np.std(df[feature])
199             mean = np.mean(df[feature])
200             global_means[feature] = round(mean,2)
201             max = df[feature].max()
202             min = df[feature].min()
203             KMeans_df[feature] = df[feature].apply(lambda x: (x-mean)/(max-min))
204             #KMeans_z_df[feature] = df[feature].apply(lambda x: 100*(x-mean)/(sigma))
205
206         #Clustering
207         X = KMeans_df.to_numpy()
208         #X_z = KMeans_z_df
209         K_list = [4,7]
210         n_iterations = 10
211         best_iterations = {}
212
213         for K in K_list:
214             kmeans = KMeans(n_clusters=K,init='random', n_init=10)
215             inertia = []
216             centers = []
217             labels = []
218             index_best = 0
219             avg_rel_errs = []
220             for i in range(n_iterations):
221                 kmeans.fit(X)
222                 inertia.append(kmeans.inertia_)
223                 centers.append(kmeans.cluster_centers_)
224                 labels.append(kmeans.labels_)
225                 if inertia[i]<inertia[index_best]:
226                     index_best = i

```

```

227
228     #Plot inertia
229     plt.plot(range(10), inertia, marker='o')
230     plt.xlabel('Iterations')
231     plt.ylabel('Inertia')
232     plt.title('Inertia over different iterations for K={}'.format(K))
233     plt.show()
234     #plot clusters
235     #Age for x axis
236     x_feat = 0
237     #Number of wins for y axis
238     y_feat = 4
239     #Normalised centers
240     centers_ = centers[index_best]
241     #Real data
242     X_real = df[features].to_numpy()
243     #Real centers
244     centers_real = np.empty(shape=(K,len(features)))
245     for j, feature in enumerate(features):
246         mean = np.mean(df[feature])
247         max = df[feature].max()
248         min = df[feature].min()
249         for i in range(K):
250             centers_real[i,j] = centers_[i,j]*(max-min)+mean
251
252     #Visualisation
253     #Normalised data
254     #plt.scatter(X[:,x_feat],X[:,y_feat],c=labels[index_best],cmap='viridis',
255     s=50,alpha=0.8)
256     #plt.scatter(centers_[:,x_feat],centers_[:,y_feat],marker='X',color='red',
257     s=200,label='Centers')
258     #plt.xlabel('age')
259     #plt.ylabel('wins')
260     #plt.legend()
261     #plt.title('KMeans clustering, K={}'.format(K))
262     #plt.show()
263
264     #Real data
265     plt.scatter(X_real[:,x_feat],X_real[:,y_feat],c=labels[index_best],
266     cmap='viridis',s=50,alpha=0.8)
267     plt.scatter(centers_real[:,x_feat],centers_real[:,y_feat],marker='X',
268     color='red',s=200,label='Centers')
269     plt.xlabel('age')
270     plt.ylabel('wins')
271     plt.legend()
272     plt.title('KMeans clustering, K={}'.format(K))

```

```

273         plt.show()
274
275         #Compare Cluster means
276         #Get separate clusters
277         print('Global means:')
278         print(global_means)
279         for j in range(K):
280             means = {}
281             for feature in features:
282                 #Standardization: z-scoring
283                 mean = np.mean(df[feature])
284                 max = df[feature].max()
285                 min = df[feature].min()
286                 means[feature] = round(np.mean(df[feature].iloc
287                     [labels[index_best] == j]),2)
288
289             cl_size = len(df.iloc[labels[index_best] == j])
290             print('Means for cluster number {} (size {}) for K={}:'.format(
291                 j+1,cl_size,K))
292             print(means)
293             relative_err = [round(100*abs(means[feature]-global_means[feature])/
294                 abs(means[feature])),2) for feature in features]
295
296             print('Relative differences to global means (%): ')
297             print(relative_err)
298             print('Average (%): ', round(sum(relative_err)/len(relative_err),2))
299             avg_rel_errs.append(sum(relative_err)/len(relative_err))
300         #Choose clusters for Bootstrap
301         if (K==4):
302             global cl1_indices
303             cl1_indices = labels[index_best] == np.argsort(avg_rel_errs)[0]
304             global cl2_indices
305             cl2_indices = labels[index_best] == np.argsort(avg_rel_errs)[1]
306     except Exception as e:
307         print(r'Unable to read data: {}'.format(str(e)))
308
309
310 def ContingencyTable(df):
311     try:
312         print('\n\nContingency table:')
313         features = ['won', 'won_diff']
314         N_Bins = 4
315         N = len(df)
316         Kbins_discret = KBinsDiscretizer(n_bins=N_Bins,encode='ordinal',
317             strategy='quantile')
318         binned_data = Kbins_discret.fit_transform(df[features])

```

```

319 df_binned = pd.DataFrame(binned_data, columns=features)
320 #Print bins interval ranges
321 for feature in features:
322     bins_indices = [df_binned[df_binned[feature]==x].index.tolist()
323                     for x in [0.0,1.0,2.0,3.0]]
324     print('For feature {}, we have:'.format(feature))
325     for i, indices in enumerate(bins_indices):
326         print('Bin number {}: [{} , {}]'.format(i+1,
327         df[feature].iloc[indices].min(),df[feature].iloc[indices].max()))
328     print('\n')
329 c_table = pd.crosstab(df_binned[features[0]], df_binned[features[1]])
330 print(c_table)
331 X = c_table.to_numpy()
332 CT1, CT2 = [], []
333 for i in range(N_Bins):
334     CT1.append(np.sum(X[:,i]))
335     CT2.append(np.sum(X[i,:]))
336 #Sum over columns
337 sum_cols = np.array(CT1)
338 #Sum over rows
339 sum_rows = np.array(CT2)
340 SumRowsMat = np.transpose(np.array([CT2,CT2,CT2,CT2]))
341 #print(SumRowsMat)
342
343 #Conditional probability
344 print('Conditional frequency table:')
345 cp_table = np.divide(X,SumRowsMat)
346 #cp_table = np.divide(X,N)
347 print(np.round(cp_table, decimals=2))
348 #Quetelet index table and Pearson's Chi-squared
349 print('Quetelet index table:')
350 q_mat = np.empty(shape=(4,4))
351 pearson_mat = np.empty(shape=(4,4))
352 for i in range(N_Bins):
353     for j in range(N_Bins):
354         q_mat[i,j] = (N*X[i,j])/(sum_rows[i]*sum_cols[j])-1
355         pearson_mat[i,j] = (X[i,j]-sum_rows[i]*sum_cols[j]/N)**2/
356         (sum_rows[i]*sum_cols[j])
357
358 print(np.round(q_mat,2))
359 print('Average Quetelet index: ', round(np.sum(q_mat*np.divide(X,N)),2))
360 print('Chi-squared: ', round(np.sum(pearson_mat),2))
361 print('Degrees of freedom: ', 9)
362 print('Pearson Chi-squared: ', round(N*np.sum(pearson_mat),2))
363
364 except Exception as e:

```

```

365         print(r'Unable to perform contingency table calculations: {}'.format(str(e)))
366
367
368 def Bootstrap(df):
369     try:
370         print('\n\nBootstrap method:')
371         N_iter = 5000
372         N = len(df)
373         feature = 'won'
374         cl1 = df.iloc[cl1_indices]
375         cl2 = df.iloc[cl2_indices]
376         cl1_idx = cl1.index.tolist()
377         cl2_idx = cl2.index.tolist()
378         indices = np.empty(shape=(N_iter,N), dtype=int)
379         means = []
380         means_cl1 = []
381         means_cl2 = []
382         for i in range(N_iter):
383             indices[i,:] = np.random.randint(0, N, size=N, dtype=int)
384             means.append(np.mean(df[feature].iloc[indices[i,:]]))
385             L1 = []
386             L2 = []
387             for index in indices[i,:].tolist():
388                 if(cl1_indices[index]):
389                     L1.append(df[feature].iloc[index])
390                 if(cl2_indices[index]):
391                     L2.append(df[feature].iloc[index])
392
393             if(L1):
394                 means_cl1.append(sum(L1)/len(L1))
395             else:
396                 means_cl1.append(0)
397             if(L2):
398                 means_cl2.append(sum(L2)/len(L2))
399             else:
400                 means_cl2.append(0)
401
402         #Confidence intervals
403         print('Confidence interval for the grand mean:')
404         #Pivotal
405         lp = np.mean(means)-1.96*np.std(means)
406         rp = np.mean(means)+1.96*np.std(means)
407         print('Pivotal Ic={}, {}'.format(lp,rp))
408         #Non pivotal
409         sorted_means = means[:]
410         sorted_means.sort()

```

```

411     lnp = sorted_means[125]
412     rnp = sorted_means[4874]
413     print('Non-pivotal Ic={}, {}'.format(lnp,rnp))
414     #Grand mean
415     print('Grand mean: ',np.mean(df[feature]))
416
417     #Compare cl1 and cl2 means
418     #Pivotal
419     m12 = [means_cl1[i] - means_cl2[i] for i in range(N_iter)]
420     print('Comparing cl1 and cl2 means:')
421     lp = np.mean(m12)-1.96*np.std(m12)
422     rp = np.mean(m12)+1.96*np.std(m12)
423     print('Pivotal Ic={}, {}'.format(lp,rp))
424     #Non pivotal
425     sorted_means = m12[:]
426     sorted_means.sort()
427     lnp = sorted_means[125]
428     rnp = sorted_means[4874]
429     print('Non-pivotal Ic={}, {}'.format(lnp,rnp))
430
431     #Compare cl1 and grand mean
432     #Pivotal
433     m1g = [means[i] - means_cl1[i] for i in range(N_iter)]
434     print('Comparing cl1 and grand mean:')
435     lp = np.mean(m1g)-1.96*np.std(m1g)
436     rp = np.mean(m1g)+1.96*np.std(m1g)
437     print('Pivotal Ic={}, {}'.format(lp,rp))
438     #Non pivotal
439     sorted_means = m1g[:]
440     sorted_means.sort()
441     lnp = sorted_means[125]
442     rnp = sorted_means[4874]
443     print('Non-pivotal Ic={}, {}'.format(lnp,rnp))
444
445     #Compare cl1 and cl2 means
446     #Pivotal
447     m2g = [means[i] - means_cl2[i] for i in range(N_iter)]
448     print('Comparing cl2 and grand mean:')
449     lp = np.mean(m2g)-1.96*np.std(m2g)
450     rp = np.mean(m2g)+1.96*np.std(m2g)
451     print('Pivotal Ic={}, {}'.format(lp,rp))
452     #Non pivotal
453     sorted_means = m2g[:]
454     sorted_means.sort()
455     lnp = sorted_means[125]
456     rnp = sorted_means[4874]

```

```

457         print('Non-pivotal Ic=[{} , {}]'.format(lnp,rnp))
458     except Exception as e:
459         print(r'Unable to perform Bootstrap: {}'.format(str(e)))
460
461
462 if __name__=='__main__':
463     df = preprocessing()
464     correlation(df)
465     PCA(df)
466     KMeansClustering(df)
467     ContingencyTable(df)
468     Bootstrap(df)

```